

Power Outages

- See the main project notebook for instructions to be sure you satisfy the rubric!
- See Project 03 for information on the dataset.
- A few example prediction questions to pursue are listed below. However, don't limit yourself to them!
 - Predict the severity (number of customers, duration, or demand loss) of a major power outage.
 - Predict the cause of a major power outage.
 - Predict the number and/or severity of major power outages in the year 2020.
 - Predict the electricity consumption of an area.

Be careful to justify what information you would know at the "time of prediction" and train your model using only those features.

Summary of Findings

Introduction

The dataset that we are using recorded the major electricity outage that occurred in the US continent with different variables that might be related to the electricity outage and effects that brought with it. There are around 1500 cases of the electricity outage between 2000 and 2016. The main goal of our project is to predict the cause of outage based on the given dataset, which is a Classification question. Our target variable is 'CAUSE.CATEGORY' and it has already been categorized into its own unique values: 'severe weather', 'intentional attack', 'system operability disruption', 'public appeal', 'equipment failure', 'fuel supply emergency', 'islanding'.

Baseline Model

The total number of features is 5. Three of them will be nominal features. The total number of ordinal features is one. The total number of quantitative will be one. We will be using 'CLIMATE.CATEGORY', 'MONTH', 'OUTAGE.DURATION', 'CLIMATE.REGION', 'ANOMALY.LEVEL' to predict the CAUSE.CATEGORY of each outage. 'CLIMATE.CATEGORY', 'MONTH', 'CLIMATE.REGION' are nominal features. 'ANOMALY.LEVEL' will be an ordinal feature. The 'OUTAGE.DURATION' will be a quantitative feature. The classifier that we use is a decision tree classifier. We decide to use the accuracy score to be the metric of our prediction since we only want to the accuracy of our prediction on the cause of the outage and the weights for inaccuracy is the same as accuracy. The accuracy score that we got is less than 0.6. It isn't very accurate, so we decide to add more features into our prediction.

Final Model

We were wondering if the climate region is related with outage duration. So we plot it out, and it seems like the north of the country seems to have longer duration. Normal and cold climates are the major climate for these regions. Since the north region of the country seems to have longer average duration, we decide to standardize our duration by the climate region by using `StdScalerByGroup`. We also decide to treat the duration as an ordinal feature since there are several causes that are more likely to cause long duration. We realized that we have too many irrelevant features, so we dropped the Month. We are still using a decision tree classifier. We select the best parameter of the grid search `cv {'classifier__max_depth': 7, 'classifier__max_leaf_nodes': 10}`. This brings our accuracy score to about 0.65.

Fairness Evaluation

For the fairness evaluation, we want to discover whether our model is biased about different years of the data. We come up with the null hypothesis that our model is not fair that the difference between the two accuracies is not due to random chance. And the alternative hypothesis is that the differences between the accuracies is due to random choice, and our model is fair. We first find the median of the year of the outages data and we separate the data into 2 groups, one is before the median year, 2008, and the other is after the year of 2008. We tested the accuracies of the 2 groups and calculated the difference between the 2 groups as our observed difference. Then we run a permutation test to see whether our observed difference will fall into our significance level, which we set it to 0.05. Eventually, we got a p-value of 0.05 after 500 times permutation, thus we rejected our null hypothesis and concluded that our model is fair towards the two separated groups.

Code

```
In [644]: import matplotlib.pyplot as plt
import numpy as np
import os
import pandas as pd
import seaborn as sns
%matplotlib inline
%config InlineBackend.figure_format = 'retina' # Higher resolution figures
```

```
In [645]: df = pd.read_excel('outage.xlsx')
columns = np.array(df.iloc[4])
measurement = np.array(df.iloc[5])
df = df.drop(df.index[[0,1,2,3,4,5]]).set_axis(columns, axis=1, inplace=False)
x = df['OUTAGE.START.DATE'].fillna('x').astype(str).str.replace('00:00:00','').replace('x','1770-1-
y = df['OUTAGE.START.TIME'].fillna('x').astype(str).str.replace('00:00:00','').replace('x','01:01:0
df['OUTAGE.START'] = pd.to_datetime(x + ' ' + y)
z = df['OUTAGE.RESTORATION.DATE'].fillna('x').astype(str).str.replace('00:00:00','').replace('x','1
h = df['OUTAGE.RESTORATION.TIME'].fillna('x').astype(str).str.replace('00:00:00','').replace('x','0
df['OUTAGE.RESTORATION'] = pd.to_datetime(z + ' ' + h)
df = df.drop(columns = ['OUTAGE.START.DATE', 'OUTAGE.START.TIME', 'OUTAGE.RESTORATION.DATE', 'OUTAGE.R
df['OUTAGE.START'] = df['OUTAGE.START'].replace(pd.to_datetime('1770-01-01 01:01:01'),np.nan)
df['OUTAGE.RESTORATION'] = df['OUTAGE.RESTORATION'].replace(pd.to_datetime('1770-01-01 01:01:01'),n
df = df.reset_index().drop(columns = ['index'])
df['OUTAGE.DURATION'] = df['OUTAGE.DURATION'].agg(lambda x : x.replace(0,np.nan))
df = df.drop('variables', axis=1).set_index('OBS')
df = df.dropna(subset = ['OUTAGE.DURATION'])
outage = df.copy()
outage
```

	YEAR	MONTH	U.S.STATE	POSTAL.CODE	NERC.REGION	CLIMATE.REGION	ANOMALY.LEVEL	CLIMATE.CATEGORY
OBS								
1	2011	7	Minnesota	MN	MRO	East North Central	-0.3	normal
2	2014	5	Minnesota	MN	MRO	East North Central	-0.1	normal
3	2010	10	Minnesota	MN	MRO	East North Central	-1.5	cold
4	2012	6	Minnesota	MN	MRO	East North Central	-0.1	normal
5	2015	7	Minnesota	MN	MRO	East North Central	1.2	warm
...
1526	2011	6	Idaho	ID	WECC	Northwest	-0.3	normal
1529	2016	7	Idaho	ID	WECC	Northwest	-0.3	normal
1530	2011	12	North Dakota	ND	MRO	West North Central	-0.9	cold
1532	2009	8	South Dakota	SD	RFC	West North Central	0.5	warm
1533	2009	8	South Dakota	SD	MRO	West North Central	0.5	warm
1398 rows × 53 columns								

```
In [646]: outage
          # checking all cols in the dataframe and determine
          # which is relevant to the number of customers affected by the outages
          outage.dtypes

YEAR                object
MONTH              object
U.S._STATE          object
POSTAL.CODE         object
NERC.REGION         object
CLIMATE.REGION      object
ANOMALY.LEVEL       object
CLIMATE.CATEGORY    object
CAUSE.CATEGORY      object
CAUSE.CATEGORY.DETAIL object
HURRICANE.NAMES     object
OUTAGE.DURATION     float64
DEMAND.LOSS.MW      object
CUSTOMERS.AFFECTED  object
RES.PRICE           object
COM.PRICE           object
IND.PRICE           object
TOTAL.PRICE         object
RES.SALES           object
COM.SALES           object
IND.SALES           object
TOTAL.SALES         object
RES.PERCEN          object
COM.PERCEN          object
IND.PERCEN          object
RES.CUSTOMERS       object
COM.CUSTOMERS       object
IND.CUSTOMERS       object
TOTAL.CUSTOMERS     object
RES.CUST.PCT        object
COM.CUST.PCT        object
IND.CUST.PCT        object
PC.REALGSP.STATE    object
PC.REALGSP.USA      object
PC.REALGSP.REL      object
PC.REALGSP.CHANGE   object
UTIL.REALGSP        object
TOTAL.REALGSP       object
UTIL.CONTRI         object
PI.UTIL.OFUSA       object
POPULATION          object
POPPCT_URBAN        object
POPPCT_UC           object
POPDEN_URBAN        object
POPDEN_UC           object
POPDEN_RURAL        object
AREAPCT_URBAN       object
AREAPCT_UC          object
PCT_LAND            object
PCT_WATER_TOT       object
PCT_WATER_INLAND    object
OUTAGE.START        datetime64[ns]
OUTAGE.RESTORATION   datetime64[ns]
dtype: object
```

```
In [647]: # checking null values

df.isna().sum(axis = 0)
df.head()
```

	YEAR	MONTH	U.S._STATE	POSTAL.CODE	NERC.REGION	CLIMATE.REGION	ANOMALY.LEVEL	CLIMATE.CATEGORY
OBS								
1	2011	7	Minnesota	MN	MRO	East North Central	-0.3	normal
2	2014	5	Minnesota	MN	MRO	East North Central	-0.1	normal
3	2010	10	Minnesota	MN	MRO	East North Central	-1.5	cold
4	2012	6	Minnesota	MN	MRO	East North Central	-0.1	normal
5	2015	7	Minnesota	MN	MRO	East North Central	1.2	warm

```
5 rows x 53 columns
```


```
In [648]: durations = outage['OUTAGE.DURATION'].value_counts()
durations
```

1.0	97
2880.0	15
300.0	14
60.0	14
1440.0	13
..	
565.0	1
7298.0	1
11700.0	1
7987.0	1
1548.0	1


```
Name: OUTAGE.DURATION, Length: 851, dtype: int64
```

```
In [649]: sns.histplot(outage['OUTAGE.DURATION'])
```

```
<AxesSubplot:xlabel='OUTAGE.DURATION', ylabel='Count'>
```



```
<AxesSubplot:xlabel='CAUSE.CATEGORY', ylabel='OUTAGE.DURATION'>
```

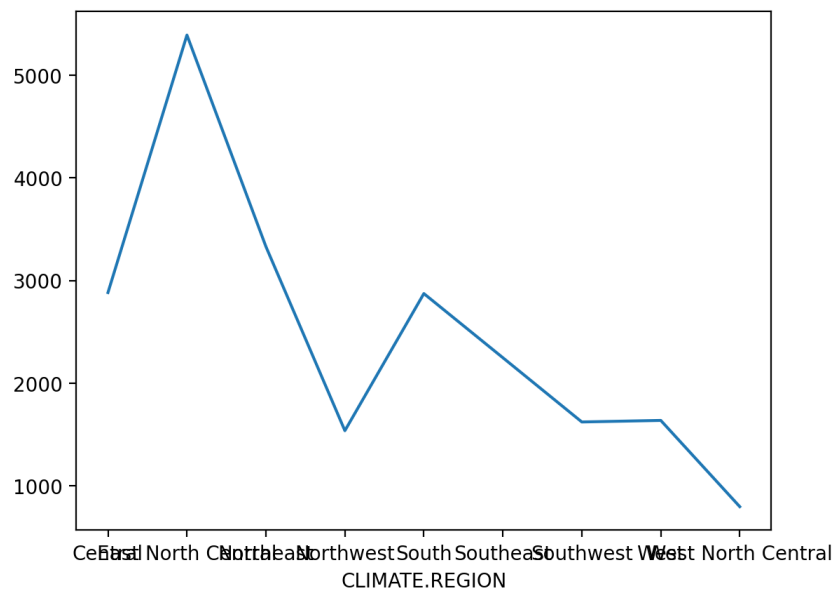


```
# duration_weather_cat
```

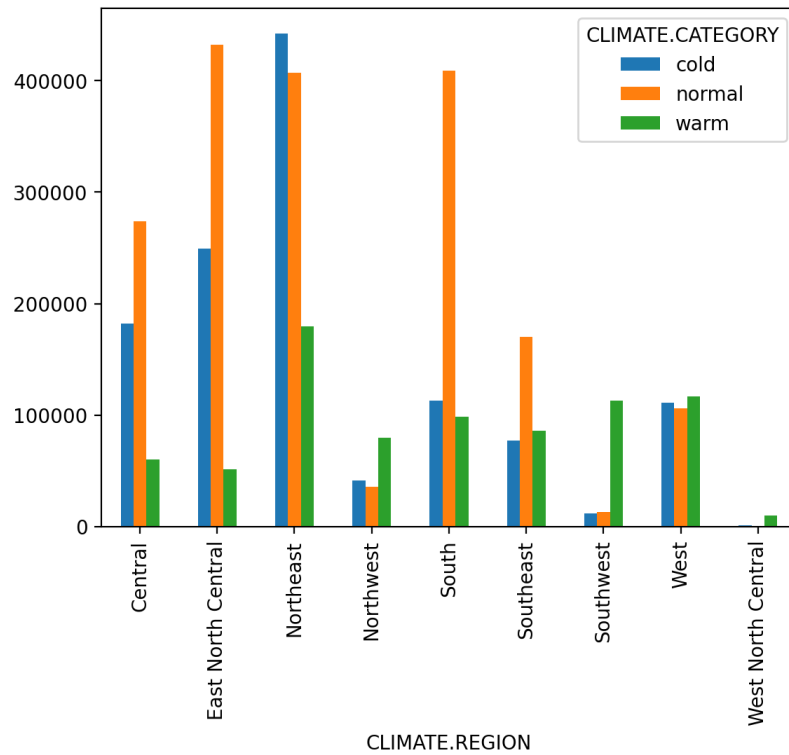
```
# outage['CAUSE.CATEGORY'].value_counts()
```

```
In [653]: df.groupby('CLIMATE.REGION')['OUTAGE.DURATION'].mean().plot()
```

<AxesSubplot:xlabel='CLIMATE.REGION'>



```
In [654]: pi = df.pivot_table(  
    index = 'CLIMATE.REGION',  
    columns = 'CLIMATE.CATEGORY',  
    values = 'OUTAGE.DURATION',  
    aggfunc = 'sum'  
)  
  
pi.plot(kind = 'bar')  
  
<AxesSubplot: xlabel='CLIMATE.REGION'>
```




```
In [655]: outage = outage[outage['OUTAGE.DURATION'].notna()]
outage.head()
```

	YEAR	MONTH	U.S._STATE	POSTAL.CODE	NERC.REGION	CLIMATE.REGION	ANOMALY.LEVEL	CLIMATE.CATEGORY
OBS								
1	2011	7	Minnesota	MN	MRO	East North Central	-0.3	normal
2	2014	5	Minnesota	MN	MRO	East North Central	-0.1	normal
3	2010	10	Minnesota	MN	MRO	East North Central	-1.5	cold
4	2012	6	Minnesota	MN	MRO	East North Central	-0.1	normal
5	2015	7	Minnesota	MN	MRO	East North Central	1.2	warm

5 rows × 53 columns

Baseline Model

```
In [686]: from sklearn.tree import DecisionTreeRegressor
from sklearn.preprocessing import FunctionTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score, train_test_split
```

```
In [687]: outage.dropna(axis=0, inplace=True)
```

```
In [688]: ## split the data into training and testing
X = outage.loc[:, ['OUTAGE.DURATION', 'CLIMATE.REGION', 'ANOMALY.LEVEL', 'MONTH', 'CLIMATE.CATEGORY', 'Y
y = outage.loc[:, ['CAUSE.CATEGORY']]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3)
```

```
In [689]: ## transform the columns
column_transformer = ColumnTransformer(
    transformers = [
        ('std', StandardScaler(), ['OUTAGE.DURATION', 'ANOMALY.LEVEL', 'YEAR']),
        ('ohh', OneHotEncoder(), ['CLIMATE.CATEGORY', 'CLIMATE.REGION', 'MONTH']),

    ], remainder = 'passthrough'
)
```

```
In [690]: ### combine transformer and random forest classifier
pred = Pipeline([
    ('col_trans', ColumnTransformer),
    ('classifier', DecisionTreeClassifier())
])
y_pred = pred.fit(X_train, y_train).predict(X_test)
print('Accuracy', metrics.accuracy_score(y_test, y_pred))
```

Accuracy 0.6578947368421053

Final Model

```
In [692]: pred.get_params().keys()

dict_keys(['memory', 'steps', 'verbose', 'col_trans', 'classifier', 'col_trans__n_jobs', 'col_trans__remainder', 'col_trans__sparse_threshold', 'col_trans__transformer_weights', 'col_trans__transformers', 'col_trans__verbose', 'col_trans__verbose_feature_names_out', 'col_trans__std', 'col_trans__ohh', 'col_trans__std_copy', 'col_trans__std_with_mean', 'col_trans__std_with_std', 'col_trans__ohh_categories', 'col_trans__ohh_drop', 'col_trans__ohh_dtype', 'col_trans__ohh_handle_unknown', 'col_trans__ohh_sparse', 'classifier__ccp_alpha', 'classifier__class_weight', 'classifier__criterion', 'classifier__max_depth', 'classifier__max_features', 'classifier__max_leaf_nodes', 'classifier__min_impurity_decrease', 'classifier__min_samples_leaf', 'classifier__min_samples_split', 'classifier__min_weight_fraction_leaf', 'classifier__random_state', 'classifier__splitter'])
```

```
In [693]: ## for ordinal encoder
outage['OUTAGE.DURATION'].max()
ran = [[i for i in range(108654)]]
```

```

In [694]: # stdscalerbygroup

from sklearn.base import BaseEstimator, TransformerMixin
import numpy

class StdScalerByGroup(BaseEstimator, TransformerMixin):

    def __init__(self):
        pass

    def fit(self, X, y=None):
        """
        :Example:
        >>> cols = {'g': ['A', 'A', 'B', 'B'], 'c1': [1, 2, 2, 2], 'c2': [3, 1, 2, 0]}
        >>> X = pd.DataFrame(cols)
        >>> std = StdScalerByGroup().fit(X)
        >>> std.grps_ is not None
        True
        """
        # X might not be a pandas DataFrame (e.g. a np.array)
        df = pd.DataFrame(X)

        # Compute and store the means/standard-deviations for each column (e.g. 'c1' and 'c2'),
        # for each group (e.g. 'A', 'B', 'C').
        # (Our solution uses a dictionary)
        mean_values=df.groupby(df.columns[0]).mean().values
        std_values=df.groupby(df.columns[0]).std().values
        self.grps_ = {'mean': mean_values , 'std':std_values }
        return self

    def transform(self, X, y=None):
        """
        :Example:
        >>> cols = {'g': ['A', 'A', 'B', 'B'], 'c1': [1, 2, 3, 4], 'c2': [1, 2, 3, 4]}
        >>> X = pd.DataFrame(cols)
        >>> std = StdScalerByGroup().fit(X)
        >>> out = std.transform(X)
        >>> out.shape == (4, 2)
        True
        >>> np.isclose(out.abs(), 0.707107, atol=0.001).all().all()
        True
        """

        # Hint: Define a helper function here!
        df=pd.DataFrame(X)

        res=pd.DataFrame()
        first_col = df.columns[0]
        unique_group=df[df.columns[0]].unique()

        row=0
        for i in unique_group:
            col=0

            temp=df.loc[df[first_col]==i]

```

```

data=temp.copy()

for j in temp.columns[1:].tolist():
    tmean=self.grps_['mean'][row][col]
    tstd=self.grps_['std'][row][col]

    data[j]=(temp[j]-tmean)/tstd

    col=col+1

data=data.drop(columns=data.columns[0])
row = row +1

res=pd.concat([res,data],ignore_index=True)

return res

```

```

In [695]: ## column transformers
          ## make outage duration into ordinal

column_tranformer = ColumnTransformer(
    transformers = [
        ('std',StandardScaler(), ['YEAR']),
        ('sbg',StdScalerByGroup(),['CLIMATE.REGION','OUTAGE.DURATION']),
        ('ordi',OrdinalEncoder(categories = ran), ['OUTAGE.DURATION']),
        ('ohh',OneHotEncoder(),['CLIMATE.CATEGORY', 'CLIMATE.REGION'])
    ], remainder = 'passthrough'
)

```

```

In [696]: pred = Pipeline([
            ('col_trans', column_tranformer),
            ('classifier', DecisionTreeClassifier(max_depth = 2))
        ])
y_pred = pred.fit(X_train,y_train).predict(X_test)
print('Accuracy',metrics.accuracy_score(y_test, y_pred))

Accuracy 0.6602870813397129

```

```

In [697]: ## setting hyperparameters,
hyperparameters = {
    'classifier__max_depth': [3,5,7,10],
    'classifier__max_leaf_nodes' : [4,6,8,10,15,20]
#     'classifier__max_depth': [i for i in np.arange(5,1300,100)]
#     'classifier__class_weight': [{0: 1, 1: 1}, {0: 1, 1/: 5}, {0: 1, 1: 1}, {0: 1, 1: 1},None]

}

```

```
In [698]: pred.set_params(classifier__max_depth = 7)
pred.set_params(classifier__max_leaf_nodes = 10)
y_pred = pred.fit(X_train,y_train).predict(X_test)
print('Accuracy',metrics.accuracy_score(y_test, y_pred))
```

Accuracy 0.6818181818181818

Fairness Evaluation

Null Hypothesis: The differences between prediction accuracies from before 2008 and after is due to random choice.

Alternative Hypothesis: The differences between the predicted accuracies are not randomly distributed.

Significance level: 0.05

```
In [724]: # finding the median of the year in dataframe
year_median = np.median(list(outage['YEAR'].unique()))
year_median
```

2008.0

```
In [725]: # grouping the year
early_year = outage[outage['YEAR'] < year_median]
late_year = outage[outage['YEAR'] >= year_median]
print(early_year.size, late_year.size)
```

17543 56551

```
In [767]: column_tranformer = ColumnTransformer(
    transformers = [
        ('std',StandardScaler(), ['YEAR', 'OUTAGE.DURATION']),
        #('sbg',StdScalerByGroup(),['CLIMATE.REGION','OUTAGE.DURATION']),
        ('ordi',OrdinalEncoder(categories = ran), ['OUTAGE.DURATION']),
        ('ohh',OneHotEncoder(),['CLIMATE.CATEGORY', 'CLIMATE.REGION'])
    ], remainder = 'passthrough'
)
```

```

In [768]: pred = Pipeline([
            ('col_trans', column_transformer),
            ('classifier', DecisionTreeClassifier(max_depth = 2))
        ])
pred.set_params(classifier__max_depth = 7)
pred.set_params(classifier__max_leaf_nodes = 10)

Pipeline(steps=[('col_trans',
                  ColumnTransformer(remainder='passthrough',
                                     transformers=[('std', StandardScaler(),
                                                    ['YEAR', 'OUTAGE.DURATION']),
                                                    ('ordi',
                                                     OrdinalEncoder(categories=[0,
                                                                              1,
                                                                              2,
                                                                              3,
                                                                              4,
                                                                              5,
                                                                              6,
                                                                              7,
                                                                              8,
                                                                              9,
                                                                              10,
                                                                              11,
                                                                              12,
                                                                              13,
                                                                              14,
                                                                              15,
                                                                              16,
                                                                              17,
                                                                              18,
                                                                              19,
                                                                              20,
                                                                              21,
                                                                              22,
                                                                              23,
                                                                              24,
                                                                              25,
                                                                              26,
                                                                              27,
                                                                              28,
                                                                              29, ...])),
                                                    ['OUTAGE.DURATION']),
                                                    ('ohh', OneHotEncoder(),
                                                     ['CLIMATE.CATEGORY',
                                                      'CLIMATE.REGION'])])),
            ('classifier',
             DecisionTreeClassifier(max_depth=7, max_leaf_nodes=10))])

```

```

In [769]: from sklearn.tree import DecisionTreeRegressor
          from sklearn.preprocessing import FunctionTransformer
          from sklearn.preprocessing import OneHotEncoder
          from sklearn.pipeline import Pipeline
          from sklearn.compose import ColumnTransformer
          from sklearn.linear_model import LinearRegression
          from sklearn.model_selection import cross_val_score, train_test_split

```

```
In [770]: # predictions on early years
X = early_year.drop('CAUSE.CATEGORY', axis=1)
y = early_year['CAUSE.CATEGORY']
X = X.loc[:, ['OUTAGE.DURATION', 'CLIMATE.REGION', 'ANOMALY.LEVEL', 'MONTH', 'CLIMATE.CATEGORY', 'YEAR']]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25)
pred.fit(X_train, y_train)
early_predictions = pred.predict(X_test)
early = metrics.accuracy_score(y_test, early_predictions)
print("Accuracy on early years: ", early)
```

Accuracy on early years: 0.7951807228915663

```
In [771]: # predictions on late years
X = late_year.drop('CAUSE.CATEGORY', axis=1)
y = late_year['CAUSE.CATEGORY']
X = X.loc[:, ['OUTAGE.DURATION', 'CLIMATE.REGION', 'ANOMALY.LEVEL', 'MONTH', 'CLIMATE.CATEGORY', 'YEAR']]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25)
pred.fit(X_train, y_train)
late_predictions = pred.predict(X_test)
late = metrics.accuracy_score(y_test, late_predictions)
print("Accuracy on late years: ", late)
```

Accuracy on late years: 0.6666666666666666

Hypothesis Test

```
In [742]: obs = abs(early - late)
obs
```

0.1426830919182347

```
In [847]: column_transformer = ColumnTransformer(
    transformers = [
        ('std', StandardScaler(), ['YEAR']),
        #('sbg', StdScalerByGroup(), ['CLIMATE.REGION', 'OUTAGE.DURATION']),
        ('ordi', OrdinalEncoder(categories = ran), ['OUTAGE.DURATION']),
        ('ohh', OneHotEncoder(handle_unknown='ignore'), ['CLIMATE.CATEGORY', 'CLIMATE.REGION'])
    ], remainder = 'passthrough'
)
```

```

In [848]: pred = Pipeline([
    ('col_trans', column_tranformer),
    ('classifier', DecisionTreeClassifier(max_depth = 6))
])
pred.set_params(classifier__max_depth = 7)
pred.set_params(classifier__max_leaf_nodes = 10)

Pipeline(steps=[('col_trans',
                  ColumnTransformer(remainder='passthrough',
                                     transformers=[('std', StandardScaler(),
                                                    ['YEAR']),
                                                    ('ordi',
                                                     OrdinalEncoder(categories=[0,
                                                                              1,
                                                                              2,
                                                                              3,
                                                                              4,
                                                                              5,
                                                                              6,
                                                                              7,
                                                                              8,
                                                                              9,
                                                                              10,
                                                                              11,
                                                                              12,
                                                                              13,
                                                                              14,
                                                                              15,
                                                                              16,
                                                                              17,
                                                                              18,
                                                                              19,
                                                                              20,
                                                                              21,
                                                                              22,
                                                                              23,
                                                                              24,
                                                                              25,
                                                                              26,
                                                                              27,
                                                                              28,
                                                                              29, ...])),
                                                    ['OUTAGE.DURATION']),
                                                    ('ohh',
                                                     OneHotEncoder(handle_unknown='ignore'),
                                                     ['CLIMATE.CATEGORY',
                                                      'CLIMATE.REGION'])])),
              ('classifier',
               DecisionTreeClassifier(max_depth=7, max_leaf_nodes=10))])

```



```

In [860]: n_repetitions = 500

differences = []
for _ in range(n_repetitions):

    # Shuffle the year
    shuffled_year = (
        outage['YEAR']
        .sample(frac=1)
        .reset_index(drop=True)
    )

    shuffled = (
        outage
        .assign(**{'Shuffled Year': shuffled_year})
    )

    early = shuffled[shuffled['Shuffled Year'] < 2008]
    late = shuffled[shuffled['Shuffled Year'] >= 2008]

    e_idx = (early.loc[:, ['OUTAGE.DURATION', 'CLIMATE.REGION', 'ANOMALY.LEVEL',
                           'MONTH', 'CLIMATE.CATEGORY', 'YEAR']].dropna(axis=0).index
             )
    l_idx = (late.loc[:, ['OUTAGE.DURATION', 'CLIMATE.REGION', 'ANOMALY.LEVEL',
                           'MONTH', 'CLIMATE.CATEGORY', 'YEAR']].dropna(axis=0).index
             )

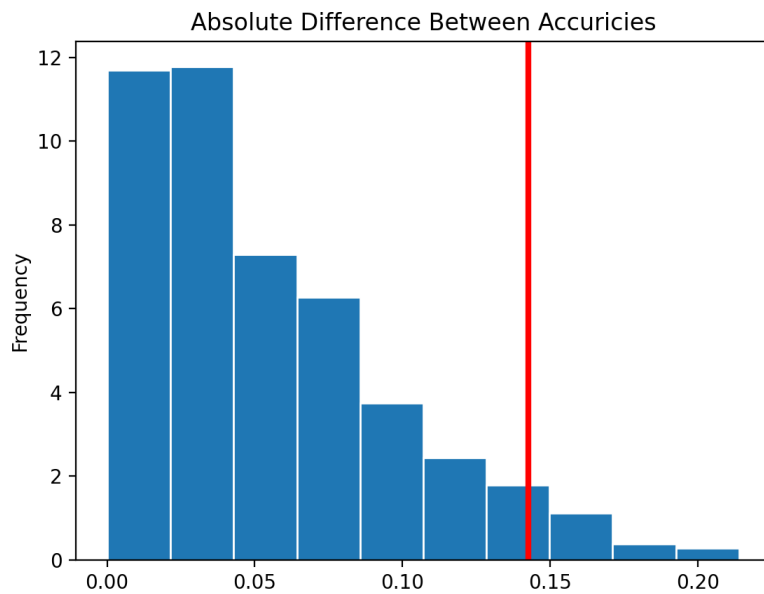
    X = early.loc[:, ['OUTAGE.DURATION', 'CLIMATE.REGION', 'ANOMALY.LEVEL', 'MONTH', 'CLIMATE.CATEGORY',
                      'e_idx = X.index
    y = early.loc[e_idx, 'CAUSE.CATEGORY']
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25)
    pred.fit(X_train, y_train)
    early_predictions = pred.predict(X_test)
    early_accuracy = metrics.accuracy_score(y_test, early_predictions)

    X = late.loc[:, ['OUTAGE.DURATION', 'CLIMATE.REGION', 'ANOMALY.LEVEL', 'MONTH', 'CLIMATE.CATEGORY',
                     'l_idx = l_idx
    y = late.loc[l_idx, 'CAUSE.CATEGORY']
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25)
    pred.fit(X_train, y_train)
    late_predictions = pred.predict(X_test)
    late_accuracy = metrics.accuracy_score(y_test, late_predictions)

    # Store the result
    diff = abs(early_accuracy - late_accuracy)
    differences.append(diff)

```

```
In [861]: title = 'Absolute Difference Between Accuracies'
pd.Series(differences).plot(kind='hist', density=True, ec='w', bins=10, title=title)
plt.axvline(x=obs, color='red', linewidth=3);
```



```
In [862]: p_val = np.mean(np.array(differences) > obs)
p_val

0.05
```

We end up rejected the null hypothesis, and concluded that the difference distribution of the accuracies between prediction of the cause categories of the outages before 2008 and the after is due to randomly choice under the 0.05 significance level. Thus, we believe our model is fair for the statistics before 2008 and after.